

Empowering RoseDB With Vector Index

Hanzhi Wu, Real Chen, Zhubo Zhou



**An efficient vector index,
when building on top of an
existing database, can be
very EASY**



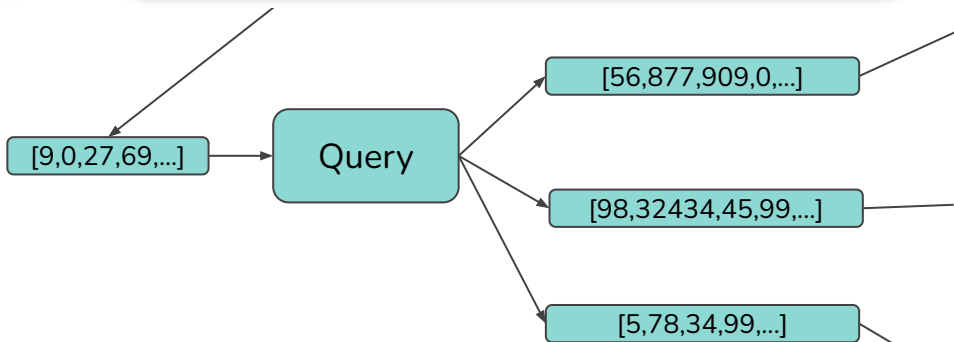
Background

🔥 Large Language Models (LLM) use **RAG** to enhance the text generation process

- Involves searching (from an external database) for **closest matching vector(s)** to the query vector produced by the LLM



how to add simple aggregate function in Velox



A row is a contiguous byte buffer. If any of the accumulators has alignment requirement, the row beginning and accumulator address will be aligned accordingly. Data is stored in the following order:

1. Null flags (1 bit per item) for
 1. Keys (only if nullable)
 2. Accumulators
2. Free flag (1 bit)
3. Keys
4. Padding for alignment
5. Accumulators, fixed-width part only
6. Variable size (32 bit)
7. Padding for alignment

```
struct AccumulatorType {  
    // Author defines data members  
    ...  
  
    // Optional. Default is true.  
    static constexpr bool is_fixed_size_ = false;  
  
    // Optional. Default is false.  
    static constexpr bool use_external_memory_ = true;  
  
    // Optional. Default is false.  
    static constexpr bool is_aligned_ = true;  
  
    explicit AccumulatorType(HashStringAllocator* allocator);  
};
```

combine


This method adds an input intermediate state to this accumulator. It receives a `HashStringAllocator*` and one `exec::arg_type<IntermediateType>` value. With default-null behavior, nulls among the input intermediate states are ignored before `combine` is called. After `combine` is called, this accumulator is assumed to be non-null.



Background

However, many existing key-value stores lack the ability to **query nearest neighbors** of a given vector efficiently

- Brute-force search: $O(n * d * \log k)$



for every query,
need to construct a
priority queue using
distances between
the given vector and
all vectors in the
database



Solution: Vector Index

- **Fast retrieval of nearest neighbors** of a given vector within a large database
- Highly **scalable** with the increase of vector dimensions
- LLM only needs answer that is good enough - Approximately closest vectors would be fine



NSW (Navigable Small Worlds) Index


GetVectors: Which is basically Dijkstra

```
C <- entry_points as min heap on distance
W <- entry_points as max heap on distance
visited <- entry_points as unordered set
while not C.empty():
    node <- pop C (nearest element in C)
    if dist(node) > dist(top W): # top W is the furthest element in W
        break
    for neighbor in neighbors of node:
        if not visited neighbor:
            visited += neighbor
            C += neighbor
            W += neighbor
            retain k-nearest elements in W
return W
```



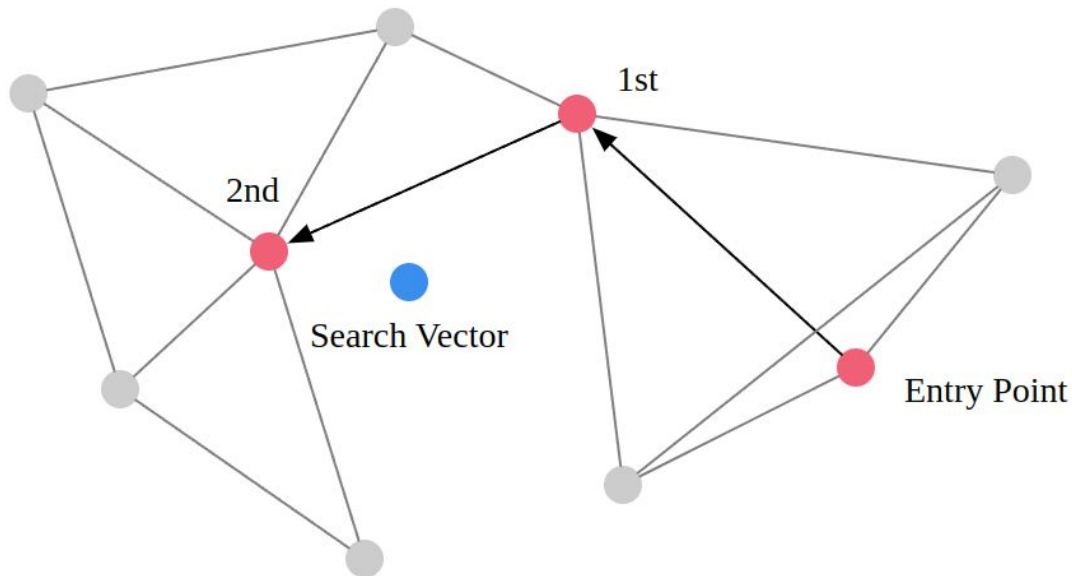


NSW (Navigable Small Worlds) Index





NSW (Navigable Small Worlds) Index






NSW (Navigable Small Worlds) Index

Put / Delete Vector: which is also EASY




NSW Index: Put Vector

Get m nearest neighbors






NSW Index: Put Vector





NSW Index: Put Vector


But graph might get too complicated ...
So cut off some edges according to **max_m**





NSW Index: Delete Vector

Delete is simply reverse the put operations





NSW Index: Parameters

Two parameters that can be configured by users

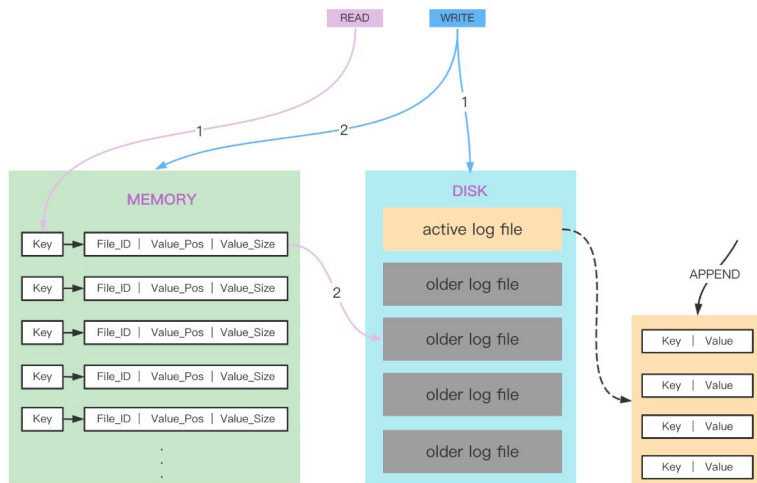
m and **max_m**



Rose DB

Integrate vector indexes to existing databases RoseDB

- A bitcask based lightweight key-value database
- Can be embedded to many existing systems







Observations From NSW Index

- Only support three operations
 - Get: get vectors “mathematically” close to the target
 - Put
 - Delete
- There is no:
 - Exact Get
 - Iterator





Integration

B+ tree index



NSW index

Tuple



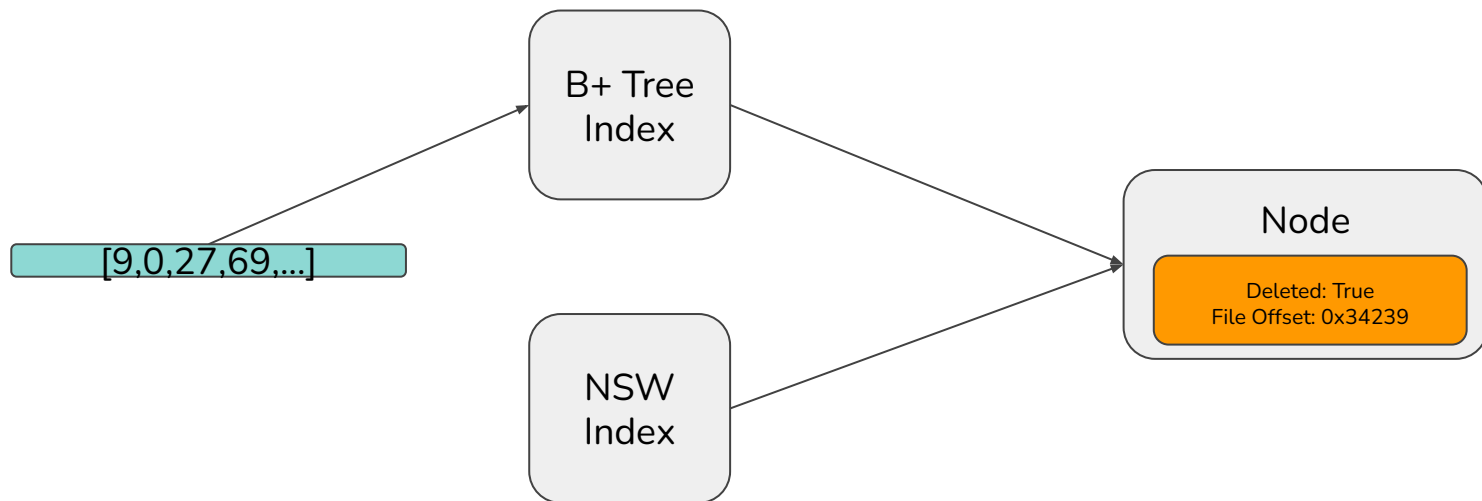


Optimization - Lazy Delete

- Observation from NSW index
 - Delete operations are super expensive
 - Taking Exclusive Latch which impact the performance
- But Delete operations in B+ tree are super fast

Optimization - Lazy Delete

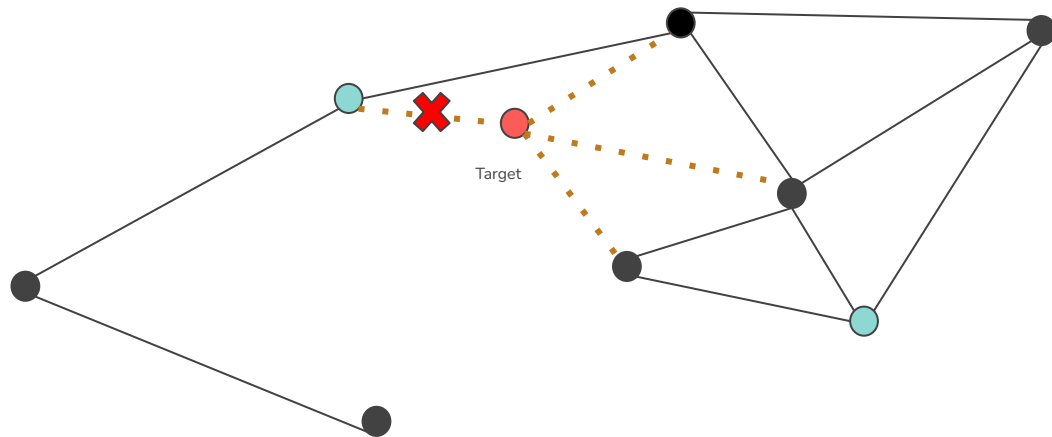
- Only delete vectors from B+ Tree
- Don't delete the vector from graph right away





Optimization - Lazy Delete


● Deleted





Optimization - Lazy Delete

Reconstruct Graph
only if 10% of nodes
are marked delete



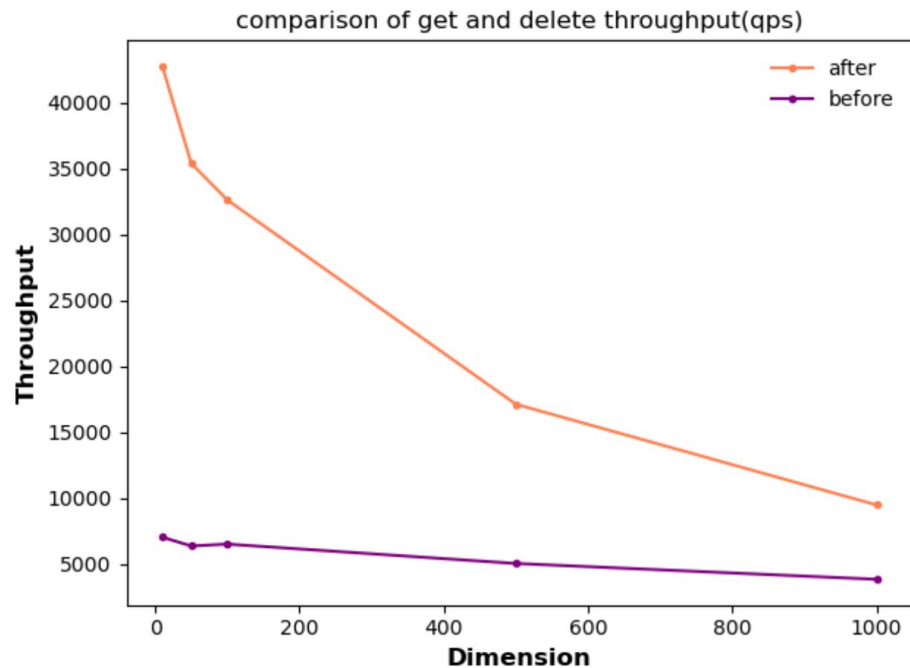


Optimization Evaluation: Setup

- Get 10000 vectors while 30% of total vectors is being deleted
 - All run **concurrently** with the following dimensions
 - 10
 - 50
 - 100
 - 500
 - 1000



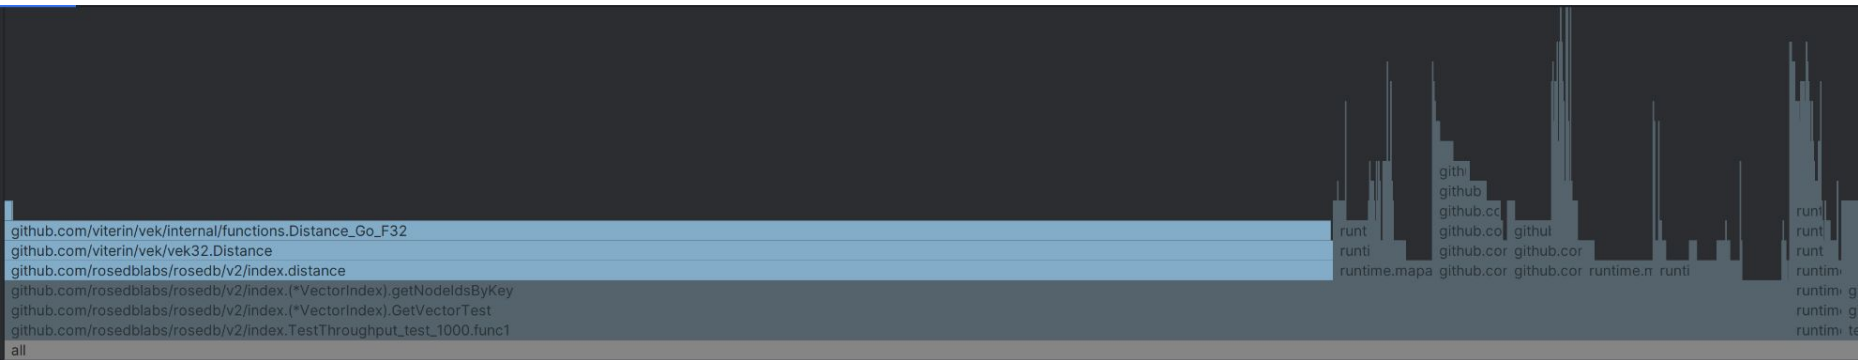
Optimization Evaluation





Optimization - Vectorized(SIMD) Execution

- In NSW index
 - Tons of vector calculations: subtract, addition, pow, sqrt...
- SIMD execution can speed up calculations





SIMD Evaluation: Environment Set Up

Arch: AMD 64 with the following SIMD register:


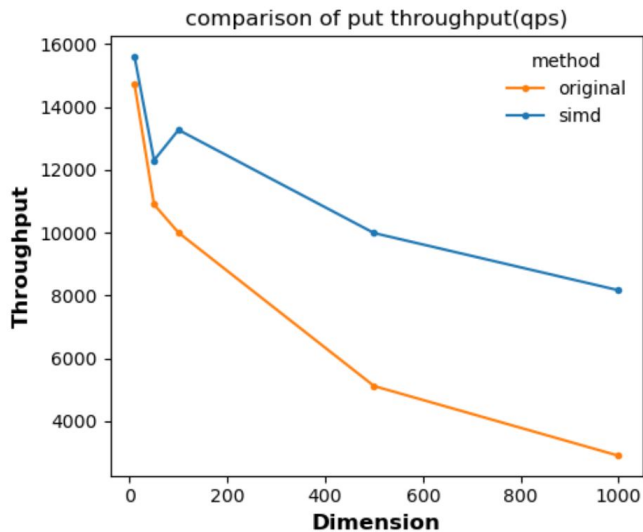
SSE2 SSE41 AVX2 BMI1 FMA SSSE3 SSE42 AES CX16 RDRAND
AVX RDSEED SSE3 PCLMULQDQ ADX POPCNT ERMS BMI2
OSXSAVE

Test Set: get 200000 times from 10000 vectors with the following dimensions:

- 10
- 50
- 100
- 500
- 1000



Optimization - Vectorized(SIMD) Execution





Lessons Learned

- NSW cannot be standalone index
 - Limited operations support
 - Bad put/delete performance
- Opens up optimization windows for integrations
 - Lazy delete
 - Lazy put could also be possible
- SIMD executions could be super helpful




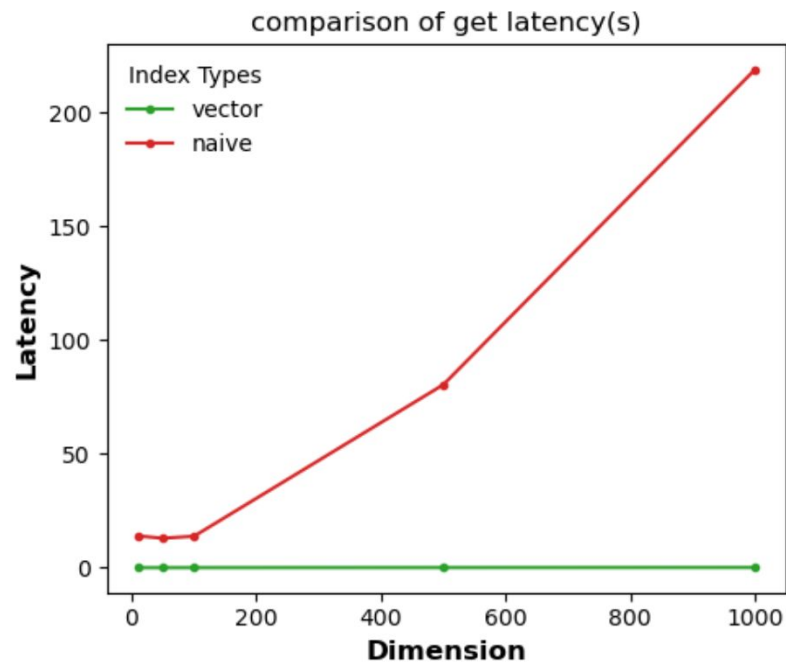
Evaluation against Naive: Environment Set Up

Arch: Apple M1 without SIMD acceleration

Test Set: get 1000 vectors from 5000 vectors with following dimensions

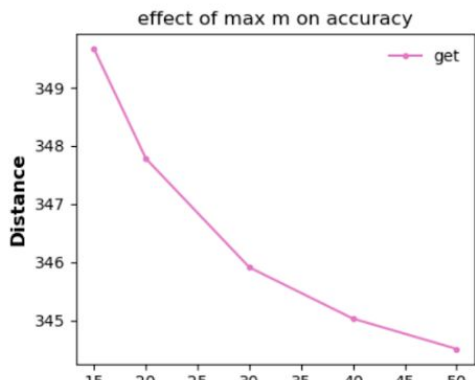
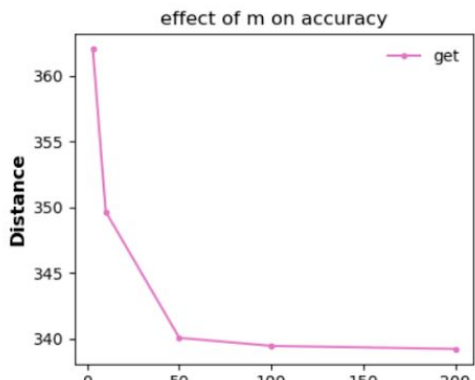
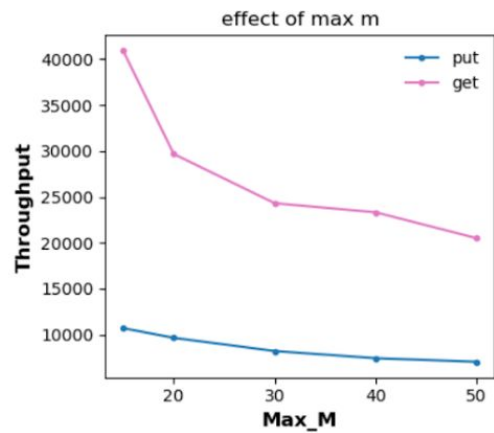

- 10
- 50
- 100
- 500
- 1000

Evaluation: Naive KNN vs NSW Index





Evaluation Against Different Parameter





Conclusion

NSW is **AWESOME!!**

**The most important
takeaway is:**



**An efficient vector index,
when building on top of an
existing database, can be
very EASY**

